



# CMPS 455

## Operating Systems

### Lecture: Virtual Memory-I

Dr. Ashok Kumar

\*Adapted from:

Operating Systems Concepts, by A. Silberschatz, P. Galvin & G.  
Gagne

Prof. John Kubiawicz, UC Berkeley  
Nachos, UC Berkeley and Univ. of Washington

# Chapter 9: Virtual Memory

---



THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*


- ❑ Background
- ❑ Demand Paging
- ❑ Copy-on-Write
- ❑ Page Replacement

- ❑ **Virtual memory** – separation of user logical memory from physical memory.
  - Only part of the program needs to be in memory for execution
  - Logical address space can therefore be much larger than physical address space
  - Allows address spaces to be shared by several processes
  - Allows for more efficient process creation
  
- ❑ Virtual memory can be implemented via:
  - Demand paging
  - Demand segmentation

# Virtual Memory That is Larger Than Physical Memory




THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*



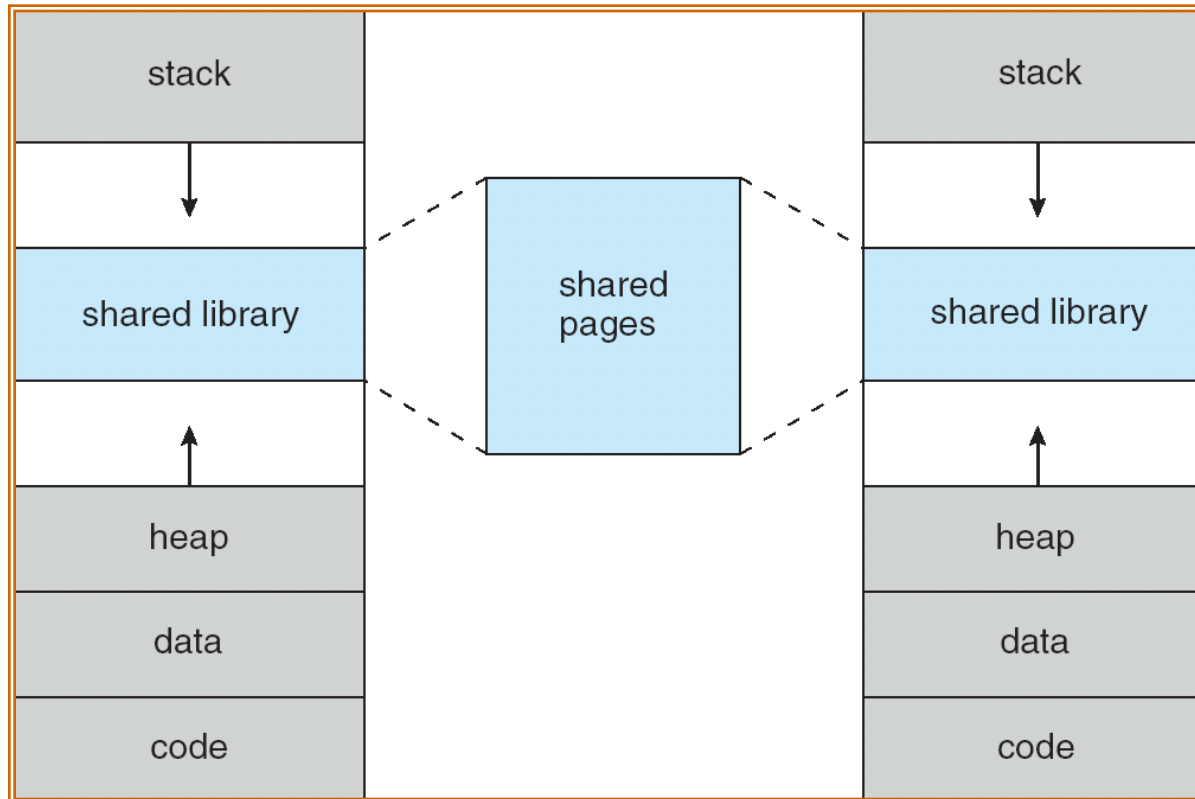
# Virtual-address Space



THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*



# Shared Library Using Virtual Memory




# Demand Paging

- ❑ Bring a page into memory only when it is needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users
  
- ❑ Page is needed  $\Rightarrow$  reference to it
  - invalid reference  $\Rightarrow$  abort
  - not-in-memory  $\Rightarrow$  bring to memory
  
- ❑ **Lazy swapper** – never swaps a page into memory unless page will be needed
  - Swapper that deals with pages is a **pager**

# Transfer of a Paged Memory to Contiguous Disk Space



THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*





# Valid-Invalid Bit



THE UNIVERSITY  
OF  
LOUISIANA  
L a f a y e t t e

- ❑ With each page table entry a valid–invalid bit is associated (**v**  $\Rightarrow$  in-memory, **i**  $\Rightarrow$  not-in-memory)
- ❑ Initially valid–invalid bit is set to **i** on all entries
- ❑ Example of a page table snapshot:

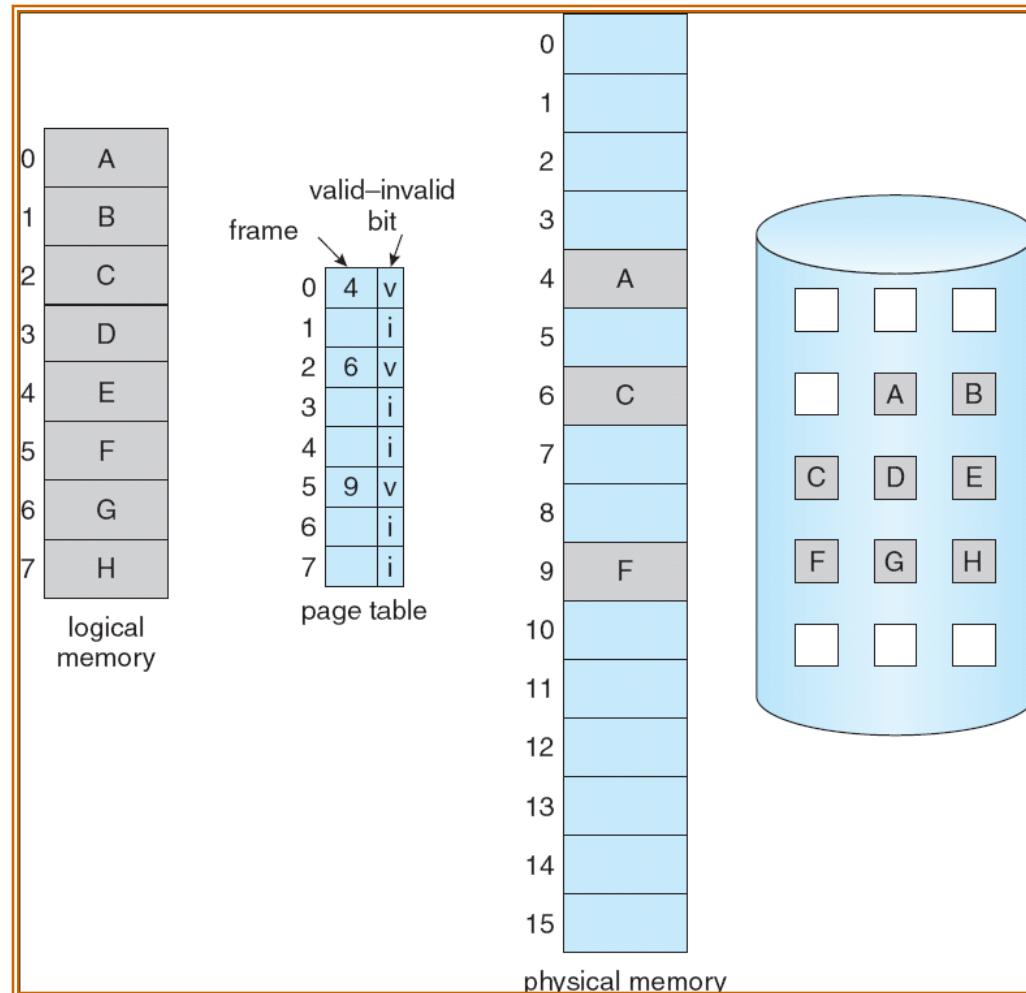
Frame #	valid-invalid bit
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>i</b>
....	
	<b>i</b>
	<b>i</b>

page table

# Page Table When Some Pages Are Not in Main Memory



THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*



- ❑ If there is a reference to a page, first reference to that page will trap to operating system:

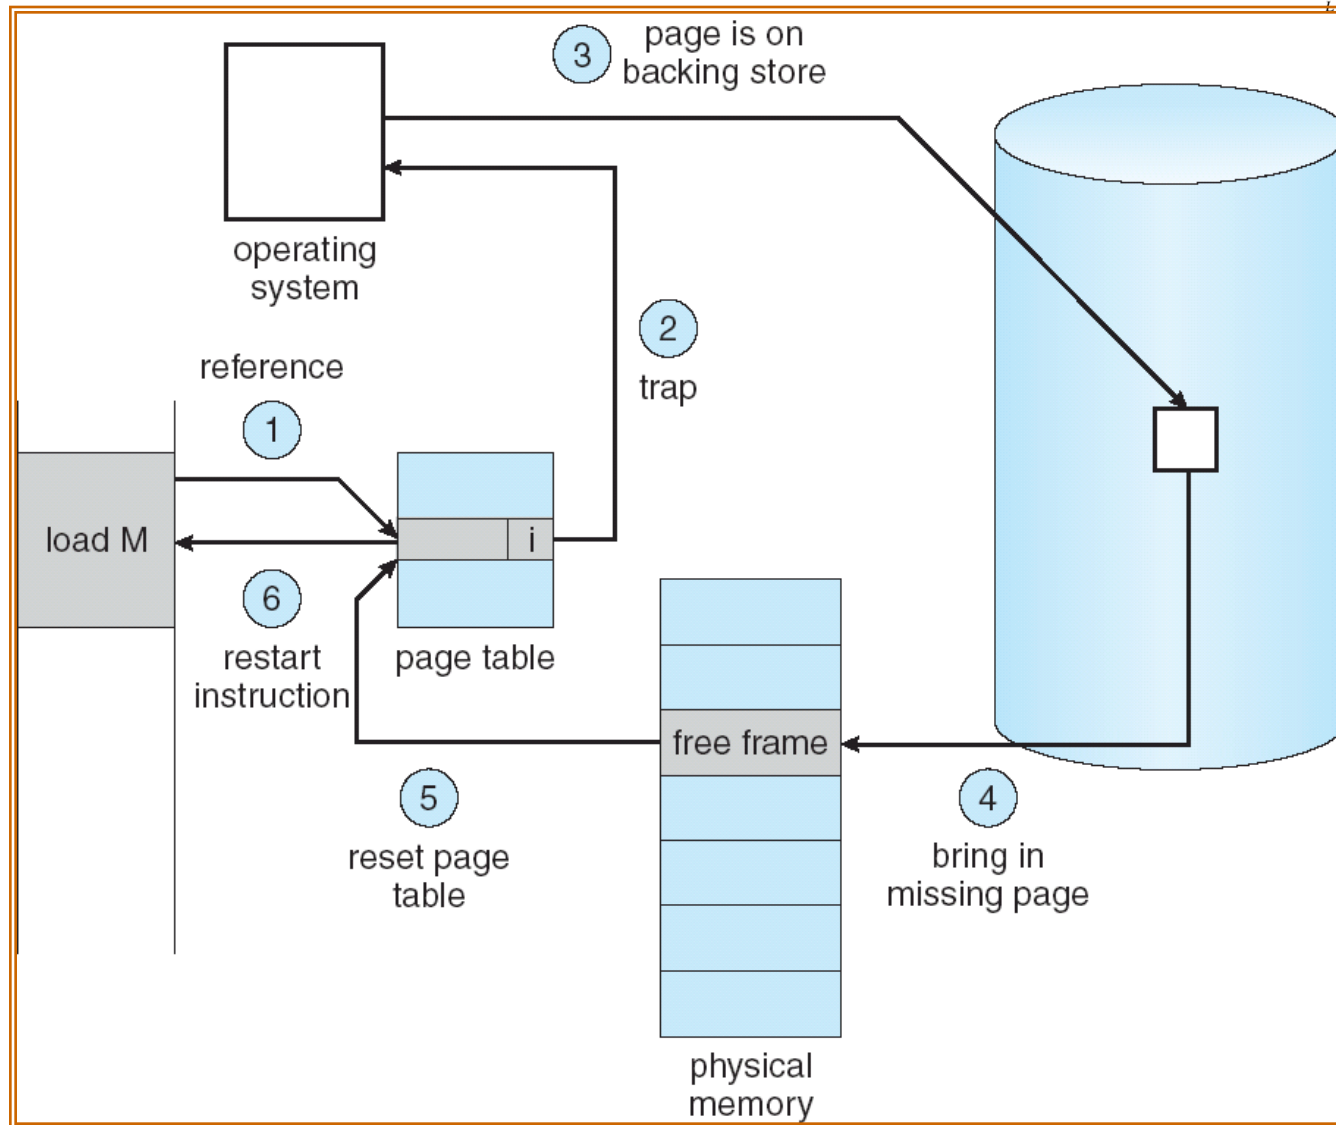
## **page fault**

1. Operating system looks at another table to decide:
  - Invalid reference  $\Rightarrow$  abort
  - Just not in memory
2. Get empty frame
3. Swap page into frame
4. Reset tables
5. Set validation bit = **v**
6. Restart the instruction that caused the page fault

# Steps in Handling a Page Fault



THE UNIVERSITY  
OF  
LOUISIANA  
at  
lafayette



# Performance of Demand Paging

- ❑ Page Fault Rate  $0 \leq p \leq 1.0$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault

- ❑ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead} \\ & ) \end{aligned}$$

# Demand Paging Example

- ❑ Memory access time = 200 nanoseconds
- ❑ Average page-fault service time = 8 milliseconds
- ❑ 
$$\begin{aligned} \text{EAT} &= (1 - p) \times 200 + p (8 \text{ milliseconds}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800 \end{aligned}$$
- ❑ If one access out of 1,000 causes a page fault, then  
$$\text{EAT} = 8.2 \text{ microseconds.}$$

This is a slowdown by a factor of 40!!

# What happens if there is no free frame?

---

- ❑ Page replacement – find some page in memory, but not really in use, swap it out
  - algorithm
  - performance – want an algorithm which will result in minimum number of page faults
- ❑ Same page may be brought into memory several times

# Page Replacement

---



THE UNIVERSITY  
OF  
LOUISIANA  
*at Lafayette*


- ❑ Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
- ❑ Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- ❑ Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory



# Need For Page Replacement



THE UNIVERSITY  
OF  
LOUISIANA  
L a f a y e t t e

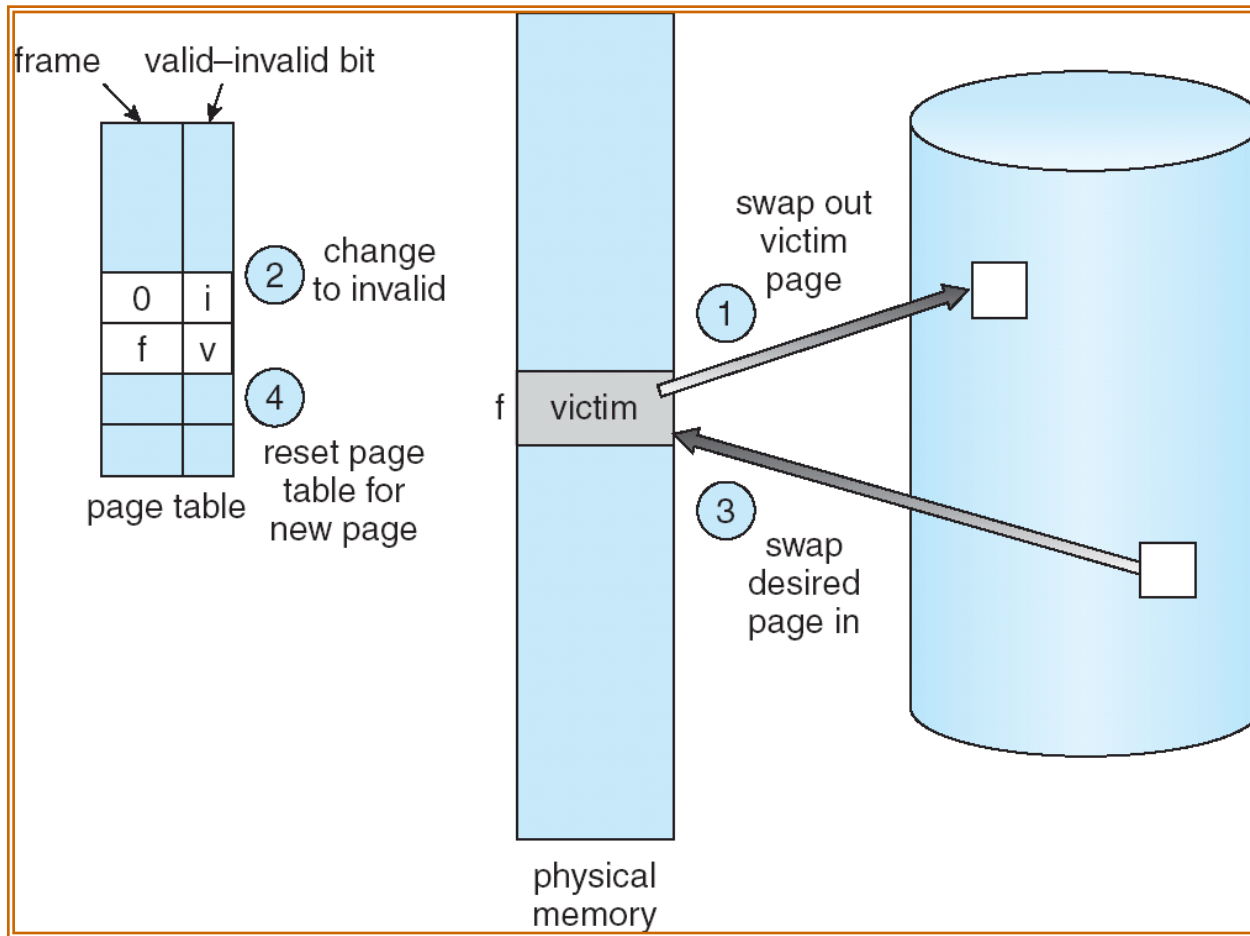


# Basic Page Replacement

---

1. Find the location of the desired page on disk
2. Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process

# Page Replacement



# Page Replacement Algorithms


- ❑ Want lowest page-fault rate
- ❑ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- ❑ In all our examples, the reference string is

**1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5**

# Graph of Page Faults Versus The Number of Frames



THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*



# First-In-First-Out (FIFO) Algorithm

- ❑ Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- ❑ 3 frames (3 pages can be in memory at a time per process)

1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

- ❑ 4 frames

1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

- ❑ Belady's Anomaly: more frames  $\Rightarrow$  more page faults

# FIFO Page Replacement



THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*


reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																	
	0	0	0																	
		1	1																	

page frames

# FIFO Illustrating Belady's Anomaly





# Optimal Algorithm



THE UNIVERSITY  
OF  
LOUISIANA  
*Leading the way to the future*

- ❑ Replace page that will not be used for longest period of time
- ❑ 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1
2
3
4

4

6 page faults

5

- ❑ How do you know this?
- ❑ Used for measuring how well your algorithm performs

# Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2			2			2				7		
	0	0	0		0		4			0			0				0		
		1	1		3		3			3			1				1		

page frames

# Least Recently Used (LRU) Algorithm

□ Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	<b>5</b>
2	2	2	2	2
3	<b>5</b>	5	<b>4</b>	4
4	4	<b>3</b>	3	3

## □ Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- When a page needs to be changed, look at the counters to determine which are to change

# LRU Page Replacement



THE UNIVERSITY  
OF  
LOUISIANA  
*L a f a y e t t e*

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames